

rna63_project_phase1

January 29, 2023

1 Project Phase One - Used Car Prices in United States

1.1 Team Background

- Project Grp 08

1.1.1 Team Members

- Team member 1
 - Name: Richard Anton
 - Email: rna63@drexel.edu

I am a MS CS student at Drexel. I expect to complete my degree in June 2023 at the end of the next quarter. I have been a software developer since 1995, and have significant experience with Python.

I have some hands on experience with machine learning, and have taken machine learning, deep learning, artificial intelligence, computer vision and DSCI-501 courses at Drexel.

I do not have as much background with exploratory data analysis. I do have a fairly solid mathematical background, including basic statistics.

A number of the prior projects that I have worked on have been focused on classification, for example of images or sentiment analysis of textual data.

Since much of my experience with handling data sets for machine learning problems has been through graduate CS courses, I have implemented them typically without a dependency on anything but NumPy and matplotlib for visualizing results, which means I have limited experience with existing common tools like scikit-learn.

I am competent at using LaTeX, but not an advanced user.

I have significant experience with NumPy and some experience using Pandas. I also have experience with matplotlib for visualizing aspects of data, but less with Seaborn.

The areas and skills I would like to grow through this project are:

1. More hands-on experience with scikit-learn and Pandas libraries.
2. More practice with visualization tools, such as matplotlib and Seaborn.
3. More practical experience with investigating a dataset, i.e. what data cleaning is needed and what relationships can be discovered in the data.
4. Experience with a practical machine learning regression problem.
5. Feature selection techniques to make machine learning problems more effective.

1.2 Topic

I would like to better understand how different variables influence the price of used vehicles. During COVID-19 supply problems, used car prices increased significantly when new car availability decreased which caused some of my curiosity in this area.

Additionally, I have always been interested in cars overall, and have bought and sold a number of older and sometimes antique vehicles, and it would be of practical use to be better able to understand what is a fair price.

This purpose is also applicable to user car vendors, consumers selling their used vehicles privately or as trade-ins, new car dealerships purchasing trade-in vehicles, and car information websites such as [edmunds.com](https://www.edmunds.com) and [Kelly Blue Book](https://www.kbb.com) that provide price range information for vehicles based on vehicle details.

I want to investigate what factors influence used car prices compared to the obvious ones of age and odometer mileage. I also want to understand other relationships between variables such as how much make and model affects the price for different age vehicles compared to the condition, or how much location affects the price of similar vehicles.

I am interested exploring a regression problem on tabular data with different types of input features since I have more experience to date with classification.

I think this is a good topic to expand my experience with visualization techniques, gain experience with practical data preparation, and practice using relevant libraries and frameworks, especially scikit-learn, Pandas, and XGBoost.

Particular techniques I want to investigate in part two of this project are

- Basic regression techniques for this problem, such as linear, ridge and lasso regression
- Regression using ensemble models, such as Random Cut Forest and XGBoost
- Feature selection techniques.
 - Using Random Forest and/or XGBoost model to evaluate feature importance.

In addition to the already described objectives, if time permits, I plan to choose the best performing regression model and embed it into a [Streamlit](https://streamlit.io/) project to create a web tool which can take a set of inputs and predict the price of a used car.

1.3 Datasets Available

There are a number of datasets available related to used cars listing or sales. They have different features, different sizes of samples and some are more prepared and preprocessed than others.

These relevant datasets located include:

- **Cargurus:** <https://www.kaggle.com/datasets/anaymital/us-used-cars-dataset>
 - This dataset contains rows with 66 columns for three million user car listings.
 - This dataset was created from Cargurus inventory in September 2020 using a crawler built by the owner of the Kaggle dataset.
- **TrueCar:** <https://www.kaggle.com/datasets/jpayne/852k-used-car-listings>
 - The provider created this dataset by scraping TrueCar.com for used car listings on 9/24/2017.
 - Columns include year, make, model, price, VIN, city, state

- **Craigslist:** <https://www.kaggle.com/datasets/austinreese/craigslist-carstrucks-data>
 - This dataset was created by the provider scraping data from Craigslist car listings.
 - The columns include price, condition, manufacturer, and latitude/longitude plus 18 other categories.
- **Carvana:** <https://www.kaggle.com/datasets/ravishah1/carvana-predict-car-prices>
 - This data is based on Carvana car sales, and only contains about 22,000 rows with four columns.
- **USDOT:** <https://catalog.data.gov/dataset/auto-sales>
 - This is available through the government bts.gov system at <https://data.bts.gov/Research-and-Statistics/Auto-Sales/7n6a-n5tz>. It seems to be primarily a historic aggregate time series of new car sales in the USA. It is part of the monthly transportation statistics published by the US Department of Transportation. It does not seem very useful to the objectives of this project. It does provide some contextual timeseries data for volume of new cars to compare against used car prices by year.

Before completing exploratory data analysis it seems likely that the Cargurus and/or Craigslist datasets would be most useful for this analysis, but this also depends on the quality of the data and effort required to prepare it within the available time so the final choice will be based on the phase one report.

1.4 Analysis of Datasets

1.4.1 Note on Datasets

Note that the code in the cells below will not run unless the data is downloaded and unzipped in advance.

The code expects a dataset directory with the following structure:

```
% find datasets -name '*.csv' -or -type d
datasets
datasets/cargurus
datasets/cargurus/used_cars_data.csv
datasets/truecar
datasets/truecar/true_car_listings.csv
datasets/craigslist
datasets/craigslist/vehicles.csv
datasets/carvana
datasets/carvana/carvana.csv
datasets/usdot
datasets/usdot/Monthly_Transportation_Statistics.csv
```

All of the CSV filenames are the original names from the ZIP filenames at the Kaggle URLs listed in the cell above.

1.4.2 Size of Datasets

```
% find . -name '*.csv' | xargs wc -l
3000599 ./datasets/cargurus/used_cars_data.csv
1233043 ./datasets/truecar/tc20171021.csv
852123 ./datasets/truecar/true_car_listings.csv
```

```
426881 ./datasets/craigslist/vehicles.csv
22001 ./datasets/carvana/carvana.csv
914 ./datasets/usdot/Monthly_Transportation_Statistics.csv
```

The graph below from the [bts.gov](https://www.bts.gov) based on data collected by the US Department of transportation shows new car sales volumen by month.



1.5 New Car Sales vs Used Cars

Since the US DOT data from [bts.gov](https://www.bts.gov) is an aggregate, monthly timeseries it will not directly help in our regression problem. This is even more true because no datasets were located that contain samples of used car prices over different points in time, rather than a point in time snapshot, so there is no way with the data we found to investigate relationships over time between new car sales and used car prices.

1.6 Exploratory Data Analysis

An analysis of the characteristics of one of the datasets is shown below using the included code cells and outputs with visualizations.

The code is structured so that the majority of it can be, and was, used to investigate and visualize the data from all four of the used car datasets. However, for reasons discussed in our summary and brevity, we only include the results for the Craigslist dataset.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from pandas.api.types import is_string_dtype
from pandas.api.types import is_numeric_dtype

import random
```

```
[2]: sns.set(style="darkgrid")
```

```
[3]: from IPython.display import display, HTML
```

```
[4]: class Heading(object):
    template = """<div style="float: left; padding: 10px;">
    <h{size}>{0}</h{size}>
    </div>"""
    def __init__(self, txt, size):
```

```

        self.txt = txt
        self.size = size

    def _repr_html_(self):
        return self.template.format(self.txt, size=self.size)

    def __repr__(self):
        return self.txt

def show_heading(txt, size=2):
    display(Heading(txt, size=size))

```

```

[5]: # This is the configuration for each dataset.
# We drop any columns which are obviously not going to be useful to use for
# ↪ data analysis or regression upon load,
# For example, VIN numbers and URLs to images.

datasets = {
    'carvana': {
        'title': "Carvana",
        'path': "datasets/carvana/carvana.csv",
        'sample': False,
        'drop_cols': []
    },
    'craigslist': {
        'title': "Craigslist",
        'path': "datasets/craigslist/vehicles.csv",
        'sample': True,
        'drop_cols':
        ↪ ['id', 'url', 'region', 'VIN', 'region_url', 'county', 'lat', 'long', 'description', 'image_url', 'po
    ],
    'cargurus': {
        'title': "Cargurus",
        'path': "datasets/cargurus/used_cars_data.csv",
        'sample': True,
        'drop_cols': ['vin', 'description', 'listing_id', 'major_options',
        ↪ 'wheel_system', 'trimId', 'sp_id', 'main_picture_url', 'latitude',
        ↪ 'longitude']
    },
    'truecar': {
        'title': "TrueCar",
        'path': "datasets/truecar/true_car_listings.csv",
        'sample': True,
        'drop_cols': ['Vin']
    }
}

```

```
[6]: # NOTE: This is where you select the dataset.
```

```
#use_dataset = 'cargurus'
use_dataset = 'craigslist'
#use_dataset = 'carvana'
#use_dataset = 'truecar'

dataset = datasets[use_dataset]
dataset_title = dataset['title']
ds_path = dataset['path']
```

```
[7]: if dataset['sample']:
    p = 0.01 # Keep 1% of the data for faster experimentation in phase 1.
    orig_df = pd.read_csv(
        ds_path,
        header=0,
        low_memory=False,
        skiprows=lambda i: i>0 and random.random() > p
    )
else:
    orig_df = pd.read_csv(ds_path, low_memory=False)
```

```
[8]: def describe_col(colname, df):
    show_heading("Column: {}".format(colname))

    if is_numeric_dtype(df[colname]):
        print("Numeric")
        if df[colname].dtype == 'int64':
            fmt = "{0:.0f}"
        else:
            fmt = "{0:.5f}"
    else:
        if is_string_dtype(df[colname]):
            print("String")
        else:
            print("Other")
        fmt = "{}"
    d = df[colname].describe().apply(fmt.format)
    display(d)
```

```
[9]: def describe_data(title, df):
    show_heading("Dataset Name: {}".format(title), size='1')
    print("Info:")
    display(df.info())
    print("Sample:")
    pd.set_option('display.max_columns', None)
    display(df.head())
```

```

print("Shape:", df.shape)
print("Types:")
display(df.dtypes)

null_ratio = df.isnull().sum()/len(df.index)
print("Null Proportion:")
display(null_ratio)

for colname in df.columns:
    describe_col(colname, df)

```

```

[10]: # Drop any columns we don't want and set df var.
df = orig_df

```

```

print("Original Columns:", df.columns)
if dataset['drop_cols']:
    df = orig_df.drop(columns=dataset['drop_cols'])
print("Using Columns:", df.columns)

```

```

Original Columns: Index(['id', 'url', 'region', 'region_url', 'price', 'year',
                        'manufacturer',
                        'model', 'condition', 'cylinders', 'fuel', 'odometer', 'title_status',
                        'transmission', 'VIN', 'drive', 'size', 'type', 'paint_color',
                        'image_url', 'description', 'county', 'state', 'lat', 'long',
                        'posting_date'],
                        dtype='object')
Using Columns: Index(['price', 'year', 'manufacturer', 'model', 'condition',
                    'cylinders',
                    'fuel', 'odometer', 'title_status', 'transmission', 'drive', 'size',
                    'type', 'paint_color', 'state'],
                    dtype='object')

```

```

[11]: # Describe the dataset.
describe_data(dataset_title, df)

```

Dataset Name: Craigslist

Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 4365 entries, 0 to 4364

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	price	4365 non-null	int64
1	year	4354 non-null	float64
2	manufacturer	4176 non-null	object
3	model	4303 non-null	object
4	condition	2623 non-null	object

```

5   cylinders      2525 non-null   object
6   fuel           4332 non-null   object
7   odometer       4311 non-null   float64
8   title_status   4290 non-null   object
9   transmission    4342 non-null   object
10  drive          3032 non-null   object
11  size           1251 non-null   object
12  type           3372 non-null   object
13  paint_color     3029 non-null   object
14  state          4365 non-null   object

```

dtypes: float64(2), int64(1), object(12)

memory usage: 511.6+ KB

None

Sample:

	price	year	manufacturer	model	condition	cylinders	\
0	6000	2007.0	mercedes-benz	e320 cdi	good	6	cylinders
1	34995	2018.0	ram	2500	NaN	NaN	
2	1200	2005.0	chevrolet	impala	fair	4	cylinders
3	27995	2012.0	ford	f250 super duty	NaN	NaN	
4	3999	2006.0	pontiac	grand prix	NaN	6	cylinders

	fuel	odometer	title_status	transmission	drive	size	type	\
0	diesel	124000.0	clean	automatic	rwd	NaN	sedan	
1	diesel	211000.0	clean	automatic	4wd	NaN	NaN	
2	gas	256806.0	clean	automatic	fwd	mid-size	sedan	
3	gas	26896.0	clean	automatic	NaN	NaN	NaN	
4	gas	207238.0	clean	automatic	fwd	NaN	sedan	

	paint_color	state
0	blue	al
1	NaN	al
2	blue	al
3	white	al
4	white	al

Shape: (4365, 15)

Types:

```

price           int64
year            float64
manufacturer     object
model            object
condition        object
cylinders        object
fuel             object
odometer         float64
title_status     object

```



```
transmission    object
drive           object
size            object
type            object
paint_color     object
state           object
dtype: object
```

Null Proportion:

```
price           0.000000
year            0.002520
manufacturer    0.043299
model           0.014204
condition       0.399084
cylinders       0.421535
fuel            0.007560
odometer        0.012371
title_status    0.017182
transmission    0.005269
drive           0.305384
size            0.713402
type            0.227491
paint_color     0.306071
state           0.000000
dtype: float64
```

Column: price

Numeric

```
count    4365
mean     17449
std      15478
min        0
25%      5790
50%     13810
75%     26880
max     139888
```

Name: price, dtype: object

Column: year

Numeric

```
count    4354.00000
mean     2011.10749
std        9.39370
min     1923.00000
25%     2008.00000
50%     2013.00000
75%     2017.00000
```

```

max      2022.00000
Name: year, dtype: object

Column: manufacturer

String

count      4176
unique       40
top        ford
freq       737
Name: manufacturer, dtype: object

Column: model

String

count      4303
unique     1960
top       f-150
freq        90
Name: model, dtype: object

Column: condition

String

count      2623
unique        6
top        good
freq     1254
Name: condition, dtype: object

Column: cylinders

String

count      2525
unique        8
top       6 cylinders
freq       961
Name: cylinders, dtype: object

Column: fuel

String

count      4332
unique        5
top        gas
freq     3622
Name: fuel, dtype: object

Column: odometer

Numeric

```

```

count      4311.00000
mean       99893.82208
std        228173.21176
min         0.00000
25%        38393.00000
50%        87224.00000
75%        136388.50000
max        999999.00000
Name: odometer, dtype: object

Column: title_status

String

count      4290
unique       6
top        clean
freq       4139
Name: title_status, dtype: object

Column: transmission

String

count      4342
unique       3
top        automatic
freq       3449
Name: transmission, dtype: object

Column: drive

String

count      3032
unique       3
top         4wd
freq       1353
Name: drive, dtype: object

Column: size

String

count      1251
unique       4
top        full-size
freq        666
Name: size, dtype: object

Column: type

String

```

```

count      3372
unique      13
top        sedan
freq       866
Name: type, dtype: object

Column: paint_color

String

count      3029
unique      12
top        white
freq       830
Name: paint_color, dtype: object

Column: state

String

count      4365
unique      51
top        ca
freq       544
Name: state, dtype: object

```

```

[12]: def show_pair_plots(df):
        show_heading("Pair Plots")
        # hue='Name',

        df = df.copy()
        # preprocess to convert booleans
        for colname in df.columns:
            if df[colname].dtype == 'bool':
                print('bool col:', colname)
                df[colname] = df[colname].replace({True: 1, False: 0})

        g = sns.pairplot(df, diag_kind='hist', height=2.5);

        plt.show()

```

```

[13]: def zscore(data):
        mean = np.mean(data)
        stdev = np.std(data)

        standardized_data = (data - mean) / stdev

        return standardized_data

```

```
[14]: def show_col_box_plots(df, rescale=True):
    show_heading("Box Plots (normalized)" if rescale else "Box Plots")
    columns_to_plot = [cn for cn in df.columns if is_numeric_dtype(df[cn])]

    fig, axes = plt.subplots(ncols=len(columns_to_plot))

    df = df.copy()
    df = df.fillna(0)

    for column, axis in zip(columns_to_plot, axes):
        data = df[column]
        if rescale:
            data = zscore(data)
        sns.boxplot(data=data, ax=axis)
        axis.set_title(column)

    plt.tight_layout()
    plt.show()

[15]: def show_col_hist(df):
    show_heading("Histograms")
    columns_to_plot = [cn for cn in df.columns if is_numeric_dtype(df[cn])]

    fig, axes = plt.subplots(ncols=len(columns_to_plot))

    for column, axis in zip(columns_to_plot, axes):
        sns.histplot(data=df[column], kde=True, ax=axis, bins=20)
        axis.set_title(column)

    plt.tight_layout()
    plt.show()

[16]: def show_bar_plots(df):
    colnames = [cn for cn in df.columns if not is_numeric_dtype(df[cn])]

    n_uniq = df[colnames].nunique()

    print("Column unique counts:")
    print([(c,n) for c,n in zip(colnames, n_uniq)])

    MAX_BAR_VALUES = 50
    keeping = [c for c,n in zip(colnames, n_uniq) if n <= MAX_BAR_VALUES]
    topn_colnames = [c for c,n in zip(colnames, n_uniq) if n > MAX_BAR_VALUES]

    counts = {c: n for c,n in zip(colnames, n_uniq)}
```

```

if len(colnames) == 0:
    return
show_heading("Bar Charts")

n_cols = 2
n_rows = (len(colnames) + n_cols - 1) // n_cols

fig = plt.figure(figsize=(12, n_rows * 7))
for idx, column in enumerate(colnames):
    axis = fig.add_subplot(n_rows, n_cols, idx + 1)

    data = df[column]

    if counts[column] > MAX_BAR_VALUES:
        g = sns.countplot(y=data, ax=axis, order=pd.
↪value_counts(df[column]).iloc[:MAX_BAR_VALUES].index)
        axis.set_title("Top {}: {}".format(MAX_BAR_VALUES, column))
    else:
        data = df[column]
        g = sns.countplot(y=data, ax=axis)
        axis.set_title(column)

plt.suptitle("Categorical Features")

plt.rcParams['figure.constrained_layout.use'] = True

plt.show()

```

```

[17]: def show_correlation(df, show_heatmap=True):
    show_heading("Correlation Heatmap")

    df = df.copy()

    for colname in df.columns:
        if is_numeric_dtype(df[colname]):
            #print("norm:", colname)
            df[colname] = zscore(df[colname])
        else:
            #print("cat:", colname)
            df[colname] = df[colname].astype("category").cat.codes

    display(df.corr(numeric_only=False))
    if show_heatmap:
        sns.heatmap(df.corr(), cmap='RdBu', vmin=-1, vmax=1) #, annot=True)
        plt.show()

```

```
[18]: show_correlation(df, show_heatmap=True)
```


Correlation Heatmap

	price	year	manufacturer	model	condition	\
price	1.000000	0.267085	-0.051286	0.029643	0.071200	
year	0.267085	1.000000	0.063659	0.101402	0.007819	
manufacturer	-0.051286	0.063659	1.000000	0.057095	-0.023646	
model	0.029643	0.101402	0.057095	1.000000	-0.018134	
condition	0.071200	0.007819	-0.023646	-0.018134	1.000000	
cylinders	-0.010712	-0.136786	-0.027465	0.003707	0.251244	
fuel	-0.078727	0.078559	0.003968	0.057213	0.168990	
odometer	-0.113833	-0.152445	0.000645	-0.011200	-0.034720	
title_status	-0.072151	-0.020098	0.003052	-0.044531	0.033827	
transmission	0.227669	0.111963	0.000405	-0.024561	0.407438	
drive	-0.079587	-0.067243	-0.040460	-0.061601	0.188861	
size	-0.170715	-0.126340	-0.018984	-0.032875	0.165844	
type	0.099585	0.139561	0.036697	-0.030548	0.179271	
paint_color	0.044649	0.062523	0.002139	0.022430	0.162128	
state	-0.014845	0.000779	-0.017804	0.025947	-0.023111	

	cylinders	fuel	odometer	title_status	transmission	\
price	-0.010712	-0.078727	-0.113833	-0.072151	0.227669	
year	-0.136786	0.078559	-0.152445	-0.020098	0.111963	
manufacturer	-0.027465	0.003968	0.000645	0.003052	0.000405	
model	0.003707	0.057213	-0.011200	-0.044531	-0.024561	
condition	0.251244	0.168990	-0.034720	0.033827	0.407438	
cylinders	1.000000	-0.102119	0.027352	0.044753	-0.057588	
fuel	-0.102119	1.000000	-0.087071	0.004226	0.237141	
odometer	0.027352	-0.087071	1.000000	0.002484	-0.076487	
title_status	0.044753	0.004226	0.002484	1.000000	-0.041567	
transmission	-0.057588	0.237141	-0.076487	-0.041567	1.000000	
drive	0.317788	0.010792	-0.006305	0.040563	0.020025	
size	0.364347	-0.068816	0.050903	0.080043	-0.206167	
type	0.198416	0.011953	-0.013708	-0.033061	0.136910	
paint_color	0.175181	0.010869	-0.034178	0.022296	0.058434	
state	-0.004031	-0.022418	0.001024	0.003091	-0.014999	

	drive	size	type	paint_color	state
price	-0.079587	-0.170715	0.099585	0.044649	-0.014845
year	-0.067243	-0.126340	0.139561	0.062523	0.000779
manufacturer	-0.040460	-0.018984	0.036697	0.002139	-0.017804
model	-0.061601	-0.032875	-0.030548	0.022430	0.025947
condition	0.188861	0.165844	0.179271	0.162128	-0.023111
cylinders	0.317788	0.364347	0.198416	0.175181	-0.004031
fuel	0.010792	-0.068816	0.011953	0.010869	-0.022418
odometer	-0.006305	0.050903	-0.013708	-0.034178	0.001024
title_status	0.040563	0.080043	-0.033061	0.022296	0.003091

transmission	0.020025	-0.206167	0.136910	0.058434	-0.014999
drive	1.000000	0.282584	0.253947	0.256204	-0.024760
size	0.282584	1.000000	0.167750	0.221036	-0.001375
type	0.253947	0.167750	1.000000	0.218585	-0.002224
paint_color	0.256204	0.221036	0.218585	1.000000	-0.040247
state	-0.024760	-0.001375	-0.002224	-0.040247	1.000000




```
[19]: def visualize_data(title, df):
      show_col_box_plots(df, rescale=False)
      show_col_box_plots(df)
      show_col_hist(df)
      show_bar_plots(df)
      show_pair_plots(df)
```

```
[20]: #visualize_data(dataset_title, df)

      show_col_box_plots(df, rescale=False)
```

Box Plots



1.6.1 Outliers

The above box plots for the numeric features without any rescaling show that there are some outliers that we need to cleanup.

```
[21]: # Cap outliers in numeric columns based on number of standard deviations.
def cap_outliers(df, z_mult):
    df = df.copy()

    for colname in df.columns:
        if is_numeric_dtype(df[colname]):
            lower_limit = df[colname].mean() - z_mult * df[colname].std()
            upper_limit = df[colname].mean() + z_mult * df[colname].std()
            print(f"Capping column: {colname} lower={lower_limit}┐
└upper={upper_limit}")

            df[colname] = np.where(
                df[colname] > upper_limit, upper_limit,
                np.where(
                    df[colname] < lower_limit, lower_limit,
                    df[colname]
```

```

    )
)

return df

```

```

[22]: show_heading("Data Cleaning")

df_cleaned = df.copy()

# Do some adjustments on particular columns based on meaning of columns.

# Drops rows without price information.
df_cleaned = df_cleaned.dropna(subset=['price'])

# No zero or negative prices.
MIN_PRICE = 1
# Keep the max price to a reasonable value.
MAX_PRICE = 300000

df_cleaned = df_cleaned.drop(df[df['price'] < MIN_PRICE].index)
df_cleaned = df_cleaned.drop(df[df['price'] > MAX_PRICE].index)

# No years in the future.
MAX_YEAR = 2023
df_cleaned['year'] = np.where(df_cleaned['year'] > MAX_YEAR, MAX_YEAR,
    ↪df_cleaned['year'])

# No huge odometer values.
MAX_MILES = 300000
df_cleaned['odometer'] = np.where(df_cleaned['odometer'] > MAX_MILES,
    ↪MAX_MILES, df_cleaned['odometer'])

# cap outliers beyond 3 std devs.
df_cleaned = cap_outliers(df_cleaned, z_mult=3)

# Impute value for missing fields.
df_cleaned['year'] = df_cleaned['year'].fillna(df_cleaned['year'].mean())
df_cleaned['odometer'] = df_cleaned['odometer'].fillna(df_cleaned['odometer'].
    ↪mean())

# TODO: replace categorical missing values with 'missing'

display(df_cleaned.describe())

```

Data Cleaning

Capping column: price lower=-26717.66351232329 upper=64619.16687136783
 Capping column: year lower=1981.9644989915205 upper=2039.757057894707

Capping column: odometer lower=-99895.95182623432 upper=287740.0281314552

	price	year	odometer
count	4019.000000	4019.000000	4019.000000
mean	18746.418432	2011.249978	93817.513564
std	14424.394534	7.687523	64001.131487
min	1.000000	1981.964499	0.000000
25%	6999.000000	2008.000000	38712.500000
50%	15300.000000	2013.000000	89078.000000
75%	27900.000000	2017.000000	137057.500000
max	64619.166871	2021.000000	287740.028131

```
[23]: describe_data("Cleaned " + dataset_title, df_cleaned)
visualize_data("Cleaned " + dataset_title, df_cleaned)
```

Dataset Name: Cleaned Craigslist

Info:

<class 'pandas.core.frame.DataFrame'>

Int64Index: 4019 entries, 0 to 4364

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	price	4019 non-null	float64
1	year	4019 non-null	float64
2	manufacturer	3840 non-null	object
3	model	3965 non-null	object
4	condition	2499 non-null	object
5	cylinders	2362 non-null	object
6	fuel	3990 non-null	object
7	odometer	4019 non-null	float64
8	title_status	3947 non-null	object
9	transmission	4003 non-null	object
10	drive	2783 non-null	object
11	size	1150 non-null	object
12	type	3104 non-null	object
13	paint_color	2808 non-null	object
14	state	4019 non-null	object

dtypes: float64(3), object(12)

memory usage: 502.4+ KB

None

Sample:

	price	year	manufacturer	model	condition	cylinders	\
0	6000.0	2007.0	mercedes-benz	e320 cdi	good	6	cylinders
1	34995.0	2018.0	ram	2500	NaN		NaN
2	1200.0	2005.0	chevrolet	impala	fair	4	cylinders
3	27995.0	2012.0	ford	f250 super duty	NaN		NaN

4	3999.0	2006.0	pontiac	grand prix	NaN	6 cylinders
---	--------	--------	---------	------------	-----	-------------

	fuel	odometer	title_status	transmission	drive	size	type	\
0	diesel	124000.0	clean	automatic	rwd	NaN	sedan	
1	diesel	211000.0	clean	automatic	4wd	NaN	NaN	
2	gas	256806.0	clean	automatic	fwd	mid-size	sedan	
3	gas	26896.0	clean	automatic	NaN	NaN	NaN	
4	gas	207238.0	clean	automatic	fwd	NaN	sedan	

	paint_color	state
0	blue	al
1	NaN	al
2	blue	al
3	white	al
4	white	al

Shape: (4019, 15)

Types:

price	float64
year	float64
manufacturer	object
model	object
condition	object
cylinders	object
fuel	object
odometer	float64
title_status	object
transmission	object
drive	object
size	object
type	object
paint_color	object
state	object
dtype:	object

Null Proportion:

price	0.000000
year	0.000000
manufacturer	0.044538
model	0.013436
condition	0.378204
cylinders	0.412292
fuel	0.007216
odometer	0.000000
title_status	0.017915
transmission	0.003981
drive	0.307539
size	0.713859

```
type          0.227669
paint_color   0.301319
state         0.000000
```

```
dtype: float64
```

```
Column: price
```

```
Numeric
```

```
count      4019.00000
mean       18746.41843
std        14424.39453
min         1.00000
25%        6999.00000
50%       15300.00000
75%       27900.00000
max       64619.16687
```

```
Name: price, dtype: object
```

```
Column: year
```

```
Numeric
```

```
count      4019.00000
mean       2011.24998
std         7.68752
min       1981.96450
25%       2008.00000
50%       2013.00000
75%       2017.00000
max       2021.00000
```

```
Name: year, dtype: object
```

```
Column: manufacturer
```

```
String
```

```
count      3840
unique       39
top         ford
freq        670
```

```
Name: manufacturer, dtype: object
```

```
Column: model
```

```
String
```

```
count      3965
unique     1861
top        f-150
freq        75
```

```
Name: model, dtype: object
```

```
Column: condition
```

String

count 2499
unique 6
top good
freq 1222

Name: condition, dtype: object

Column: cylinders

String

count 2362
unique 8
top 6 cylinders
freq 898

Name: cylinders, dtype: object

Column: fuel

String

count 3990
unique 5
top gas
freq 3344

Name: fuel, dtype: object

Column: odometer

Numeric

count 4019.00000
mean 93817.51356
std 64001.13149
min 0.00000
25% 38712.50000
50% 89078.00000
75% 137057.50000
max 287740.02813

Name: odometer, dtype: object

Column: title_status

String

count 3947
unique 6
top clean
freq 3797

Name: title_status, dtype: object

Column: transmission

String

```
count      4003
unique      3
top      automatic
freq      3131
Name: transmission, dtype: object
```

Column: drive

String

```
count      2783
unique      3
top      4wd
freq      1229
Name: drive, dtype: object
```

Column: size

String

```
count      1150
unique      4
top      full-size
freq      611
Name: size, dtype: object
```

Column: type

String

```
count      3104
unique      13
top      sedan
freq      798
Name: type, dtype: object
```

Column: paint_color

String

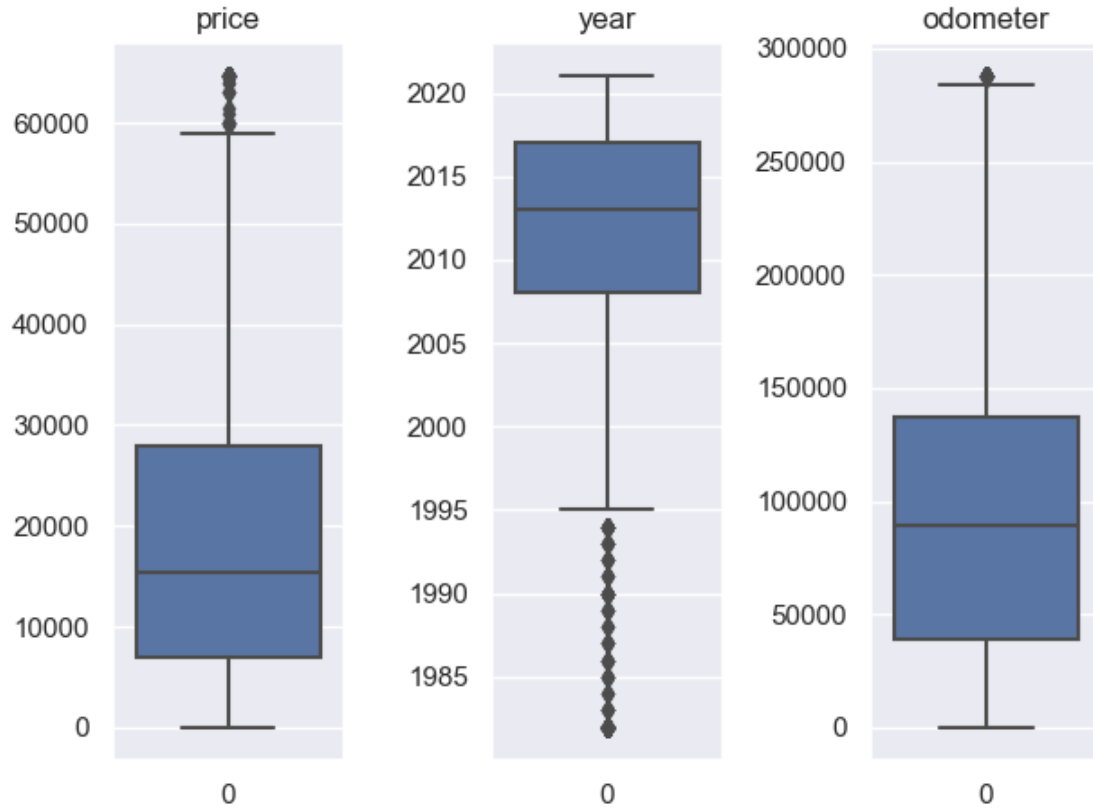
```
count      2808
unique      12
top      white
freq      761
Name: paint_color, dtype: object
```

Column: state

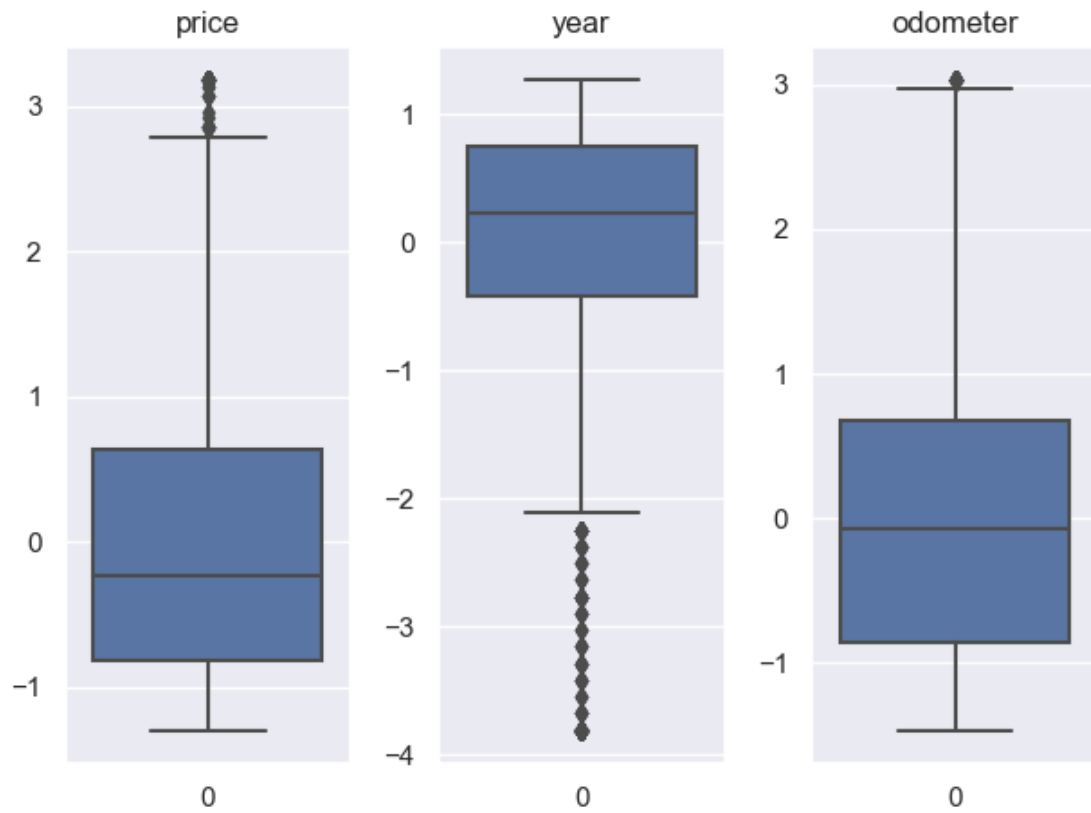
String

```
count      4019
unique      51
top      ca
freq      494
Name: state, dtype: object
```

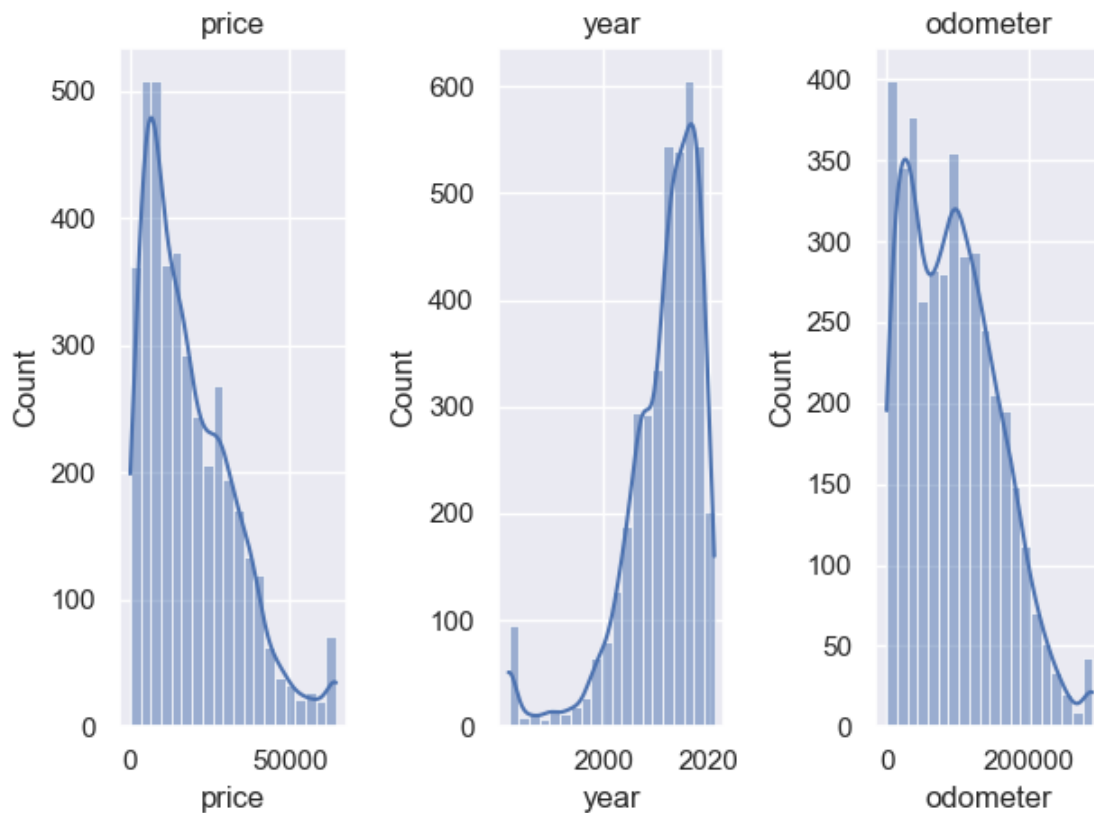
Box Plots



Box Plots (normalized)



Histograms




Column unique counts:


```
[('manufacturer', 39), ('model', 1861), ('condition', 6), ('cylinders', 8),
 ('fuel', 5), ('title_status', 6), ('transmission', 3), ('drive', 3), ('size',
 4), ('type', 13), ('paint_color', 12), ('state', 51)]
```

Bar Charts

Categorical Features



Pair Plots



1.7 Conclusions from Visualizations

The pair plots show that price is related significantly to the year and odometer values.

The histograms for price and odometer show distributions skewed towards more frequent values on the lower side, while the histogram for the year shows a distribution skewed towards more frequent higher (recent) years.

```
[24]: # Let's reexamine the correlation heatmap after cleaning up the outliers and
      ↪missing values.
      show_correlation(df, show_heatmap=True)
```


Correlation Heatmap

	price	year	manufacturer	model	condition	\
price	1.000000	0.267085	-0.051286	0.029643	0.071200	
year	0.267085	1.000000	0.063659	0.101402	0.007819	
manufacturer	-0.051286	0.063659	1.000000	0.057095	-0.023646	
model	0.029643	0.101402	0.057095	1.000000	-0.018134	
condition	0.071200	0.007819	-0.023646	-0.018134	1.000000	
cylinders	-0.010712	-0.136786	-0.027465	0.003707	0.251244	
fuel	-0.078727	0.078559	0.003968	0.057213	0.168990	
odometer	-0.113833	-0.152445	0.000645	-0.011200	-0.034720	
title_status	-0.072151	-0.020098	0.003052	-0.044531	0.033827	
transmission	0.227669	0.111963	0.000405	-0.024561	0.407438	
drive	-0.079587	-0.067243	-0.040460	-0.061601	0.188861	
size	-0.170715	-0.126340	-0.018984	-0.032875	0.165844	
type	0.099585	0.139561	0.036697	-0.030548	0.179271	
paint_color	0.044649	0.062523	0.002139	0.022430	0.162128	
state	-0.014845	0.000779	-0.017804	0.025947	-0.023111	

	cylinders	fuel	odometer	title_status	transmission	\
price	-0.010712	-0.078727	-0.113833	-0.072151	0.227669	
year	-0.136786	0.078559	-0.152445	-0.020098	0.111963	
manufacturer	-0.027465	0.003968	0.000645	0.003052	0.000405	
model	0.003707	0.057213	-0.011200	-0.044531	-0.024561	
condition	0.251244	0.168990	-0.034720	0.033827	0.407438	
cylinders	1.000000	-0.102119	0.027352	0.044753	-0.057588	
fuel	-0.102119	1.000000	-0.087071	0.004226	0.237141	
odometer	0.027352	-0.087071	1.000000	0.002484	-0.076487	
title_status	0.044753	0.004226	0.002484	1.000000	-0.041567	
transmission	-0.057588	0.237141	-0.076487	-0.041567	1.000000	
drive	0.317788	0.010792	-0.006305	0.040563	0.020025	
size	0.364347	-0.068816	0.050903	0.080043	-0.206167	
type	0.198416	0.011953	-0.013708	-0.033061	0.136910	
paint_color	0.175181	0.010869	-0.034178	0.022296	0.058434	
state	-0.004031	-0.022418	0.001024	0.003091	-0.014999	

	drive	size	type	paint_color	state
price	-0.079587	-0.170715	0.099585	0.044649	-0.014845
year	-0.067243	-0.126340	0.139561	0.062523	0.000779
manufacturer	-0.040460	-0.018984	0.036697	0.002139	-0.017804
model	-0.061601	-0.032875	-0.030548	0.022430	0.025947
condition	0.188861	0.165844	0.179271	0.162128	-0.023111
cylinders	0.317788	0.364347	0.198416	0.175181	-0.004031
fuel	0.010792	-0.068816	0.011953	0.010869	-0.022418

odometer	-0.006305	0.050903	-0.013708	-0.034178	0.001024
title_status	0.040563	0.080043	-0.033061	0.022296	0.003091
transmission	0.020025	-0.206167	0.136910	0.058434	-0.014999
drive	1.000000	0.282584	0.253947	0.256204	-0.024760
size	0.282584	1.000000	0.167750	0.221036	-0.001375
type	0.253947	0.167750	1.000000	0.218585	-0.002224
paint_color	0.256204	0.221036	0.218585	1.000000	-0.040247
state	-0.024760	-0.001375	-0.002224	-0.040247	1.000000



1.8 Findings

1.8.1 Selection of Dataset for Continued Analysis

The Carvana dataset provides very few columns(Name,Year,Miles,Price), and a much smaller set of samples than the other datasets.

The CarGurus dataset has the largest rows and the most feature columns available(at 66). However, it only provides location data as zip code, longitude and latitude, which would make analyzing the effect on price from different areas more difficult since would prefer to just compare prices across US states.

Both the Craigslist and Carvana datasets have some data cleaning work required to remove outliers as well as significant numbers of null values.

The TrueCar dataset is by far the cleanest dataset, having no null values, and more columns than the Carvana dataset. However, the columns available would only allow some of the investigations we have in mind. It provides Price, Year, Mileage, City, State, Vin, Make, Model. Vin is of little use without advanced preprocessing based on VIN coding, so the only advantage in features it provides to the Carvana dataset are from location based on City, State.

The Craigslist dataset is most amenable to the investigations planned in terms of the columns available. This is because it includes a column for the state location, the same basic make, model, price, mileage and year information common to all the datasets, but also has a number of additional feature columns which may affect price. The complete columns are: id,url,region,region_url,price,year,manufacturer,model,condition,cylinders,fuel,odometer,title_status,transmission,VIN,drive,size,type,paint_color,image_url,description,

1.8.2 Features

It is clear from the correlation values that price is correlated with year of the car. And it is also apparent that the odometer mileage is negatively correlated with the price. These are not surprising findings.

Other relationships are less obvious from the graphs. This is likely to be partly due to most of the features being categorical and not necessarily ordered in a meaningful way before correlations were calculated.

There are a significant number of null values remaining in the categorical data after the data cleaning we have already done. For example, cylinders column is missing in over 41% of the rows.

Null Value Proportion by Column

price	0.000000
year	0.002588
manufacturer	0.039755
model	0.011291
condition	0.396377
cylinders	0.417314
fuel	0.006351
odometer	0.011997
title_status	0.020936
transmission	0.006351
drive	0.312162
size	0.717008
type	0.215949
paint_color	0.310750
state	0.000000

1.9 Limitations

During this phase used sampled 1% of the larger datasets to speed up initial analysis and debugging of code. Phase two of the project will use the full Craigslist dataset.

Only preliminary data cleaning has been done. Imputation for missing values that remain is still needed.

1.10 Phase Two Plans

The data is able to support most of the investigations planned for phase two.

Specifically, for phase two we will go deeper into answering these questions:

- How much make and model, condition and other features affect prices of vehicles of the same age.
- How location affects the price of similar vehicles.
- How well a regression model built for this data will perform on test data.
- Which features are most useful for a regression model.

However, I was not able to locate datasets suitable to measuring the relationship between new car sales and used car prices.

For phase two of the project, the plan is to:

1. Begin by completing data cleanup of the Craigslist dataset, especially imputation of missing categorical feature values.
2. Continue exploratory data analysis into the price effect of different features.
3. Create a baseline linear regression model to predict price.
4. Create lasso and ridge regression models for price to see how they compare to baseline.
5. Use XGBoost library to create a regression model and also evaluate feature importance.
6. Create a report and presentation based on the findings.

If time is available to do additional work beyond the plan above, then I also plan to evaluate using Random Cut Forest for regression and evaluating feature importance and compare the results with XGBoost.

Lastly, time permitting, I plan to use [Streamlit.io](https://streamlit.io) to create a web interface for serving an inference endpoint so that a user can input the feature values for a used car and see the price the model predicts.

The report and code to reproduce it will be available via GitHub to interested parties (professor, classmates, other peers) as well as stakeholders interested in how predictable used car prices are with this model who can view the report and use the online inference web interface.